

2022 Synopsys ARC 盃 AIoT 設計應用競賽 決賽作品

非接觸式控制面板

Contactless Control Panel

指導教授：張錫嘉 教授

隊員：陳冠瑋、曹家輔、李家毓、鄭紹文

隊名：綠洲熊與他們的窩



July 22th, 2022

Agenda

- **Introduction**
- **Difficulties & Innovation**
- **Design & Implementation**
- **Result & Demo**
- **Q&A**

Agenda

- **Introduction**
- Difficulties & Innovation
- Design & Implementation
- Result & Demo
- Q&A

Introduction



**Virtual elevator button panel
& Visitor verification**




■ Face verification

- Hardware: **ARC EM9D AIoT DK**
- Model: **Siamese Network**

■ Virtual button panel

- Hardware: NVIDIA Jetson Nano



Contactless ticket machine screen



Smart door lock



Agenda

- Introduction
- **Difficulties & Innovation**
- Design & Implementation
- Result & Demo
- Q&A

Difficulties - Face Verification (ARC EM9D)

■ Memory size

- Available memory size < 2MB
- Auto-Encoder: Dimension Reduction
- **Post training quantization: fp32 → int8** ✓

■ Input images pairs for the Siamese network

- Siamese network needs to compare two images at the same time

■ Transmission time

- Validation data transmission
- **Save images for comparison in the system memory** ✓

■ MLI Library does not support some Ops

sigmoid ✗
tanh ✗




logistic ✓



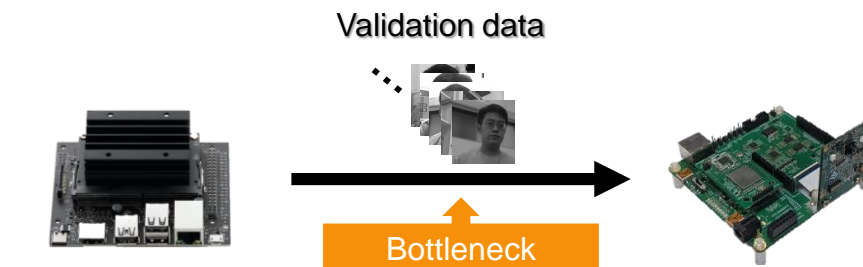
abs (int8) ✗



mul (int8) ✓



Input images pairs for Siamese network



Transmission time

Innovation - Face Verification (ARC EM9D)

■ Siamese network architecture on TinyML

- Discriminate the similarity between two different input
- Small number of samples are needed

■ Small face verification model

- **Low power:** Used in unreliable power supply scenarios
- **Less computing & memory resources:** Faster inference, faster deployment


■ Input dimension reduction

- Auto-Encoder: Reduce the dimension of the input features
- Solve the bottleneck of inference time
- Deploy a more efficient model

Innovation - Virtual Button Panel (NVIDIA Jetson Nano)

■ User - Friendly

- Click the virtual panel through the gesture
- **Index Finger**: Select the virtual button
- **Index Finger + Middle Finger**: Click the virtual button

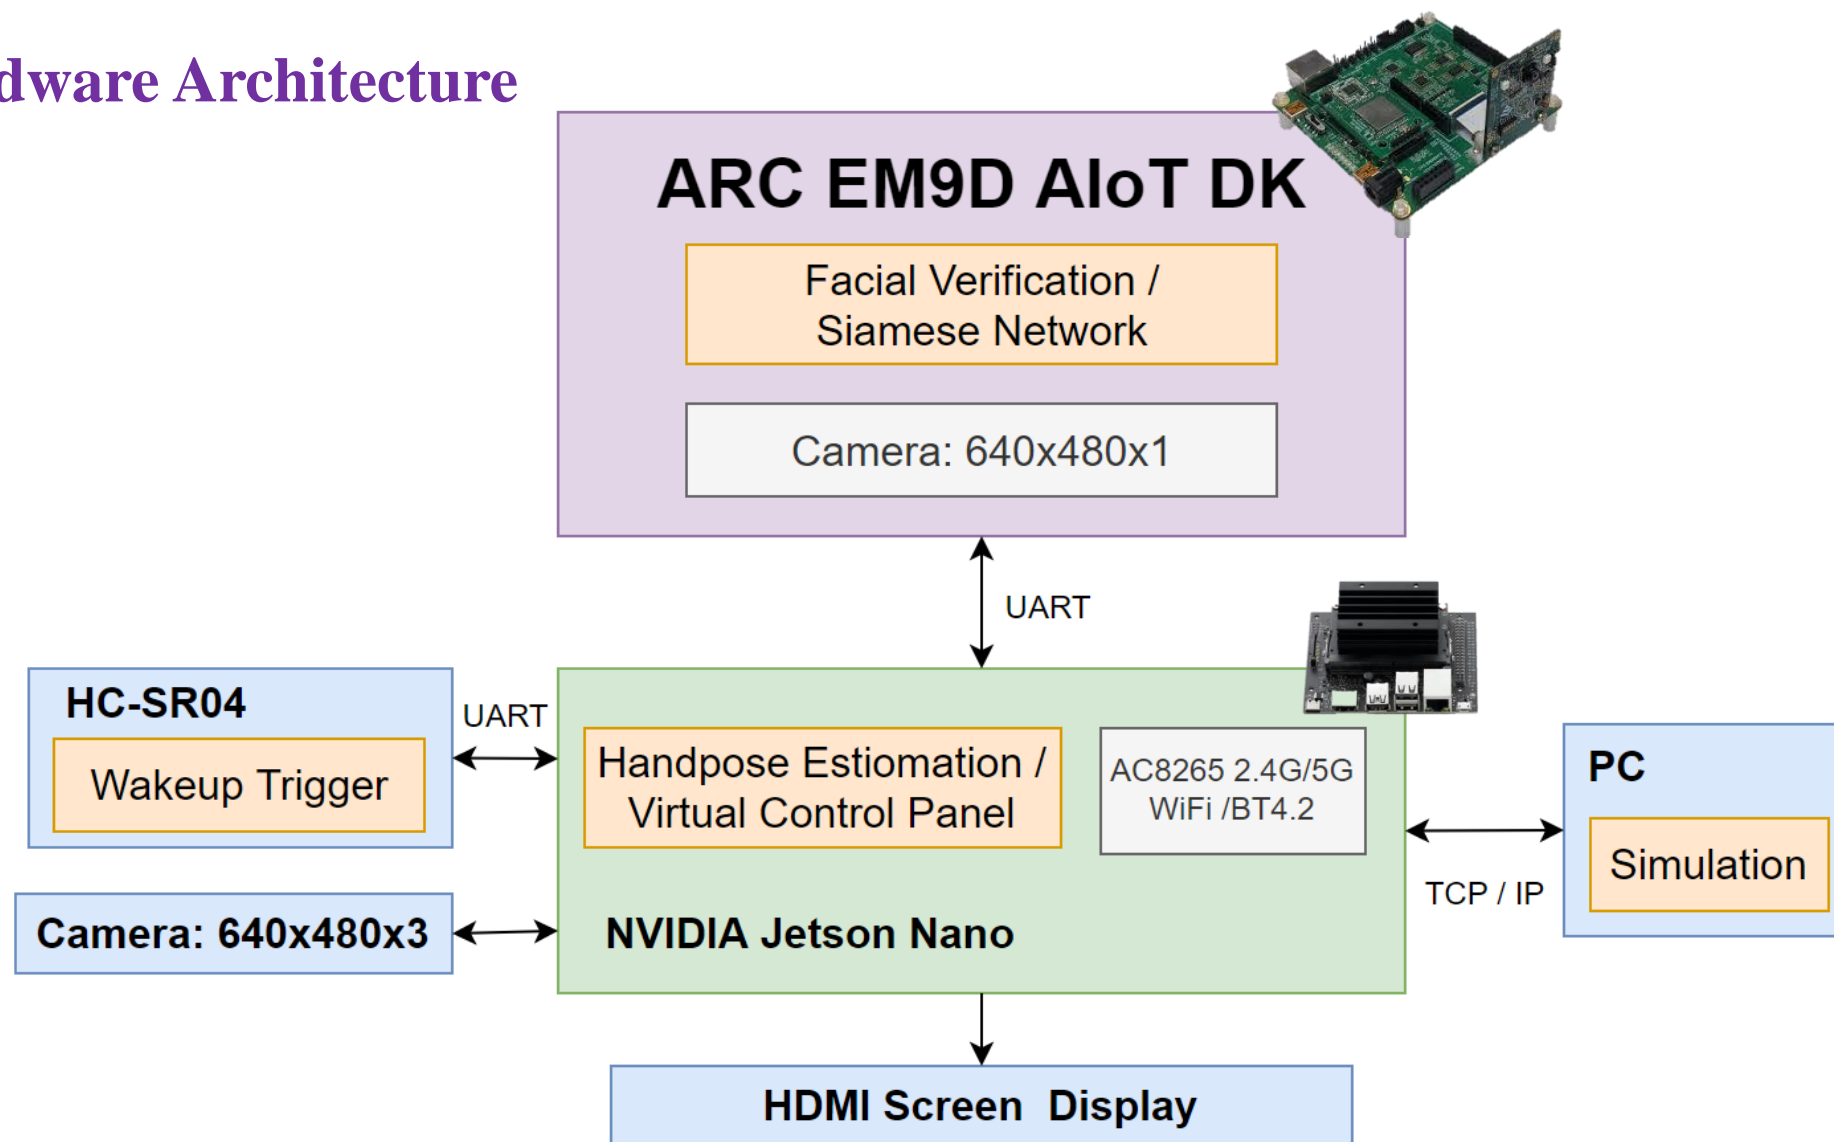


Agenda

- Introduction
- Difficulties & Innovation
- **Design & Implementation**
- Result & Demo
- Q&A

Design & Implementation


■ Hardware Architecture



Design & Implementation

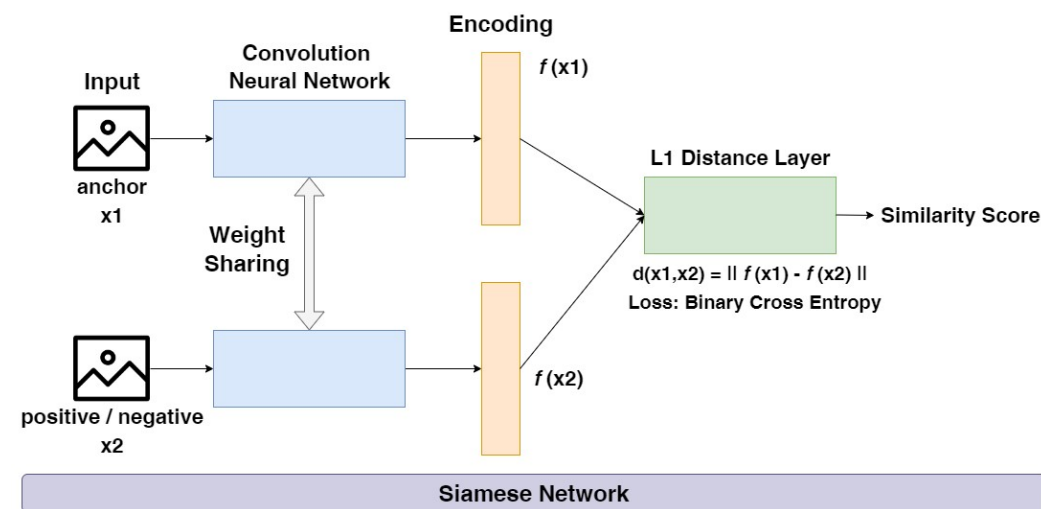

■ Siamese Network

- Inspired by “Siamese twins”
- Classify if the two inputs are the same or different
- Take two different inputs passed through two same subnetworks
- **Convert the classification problem to a similarity problem**



■ Why choose “Siamese Network”?

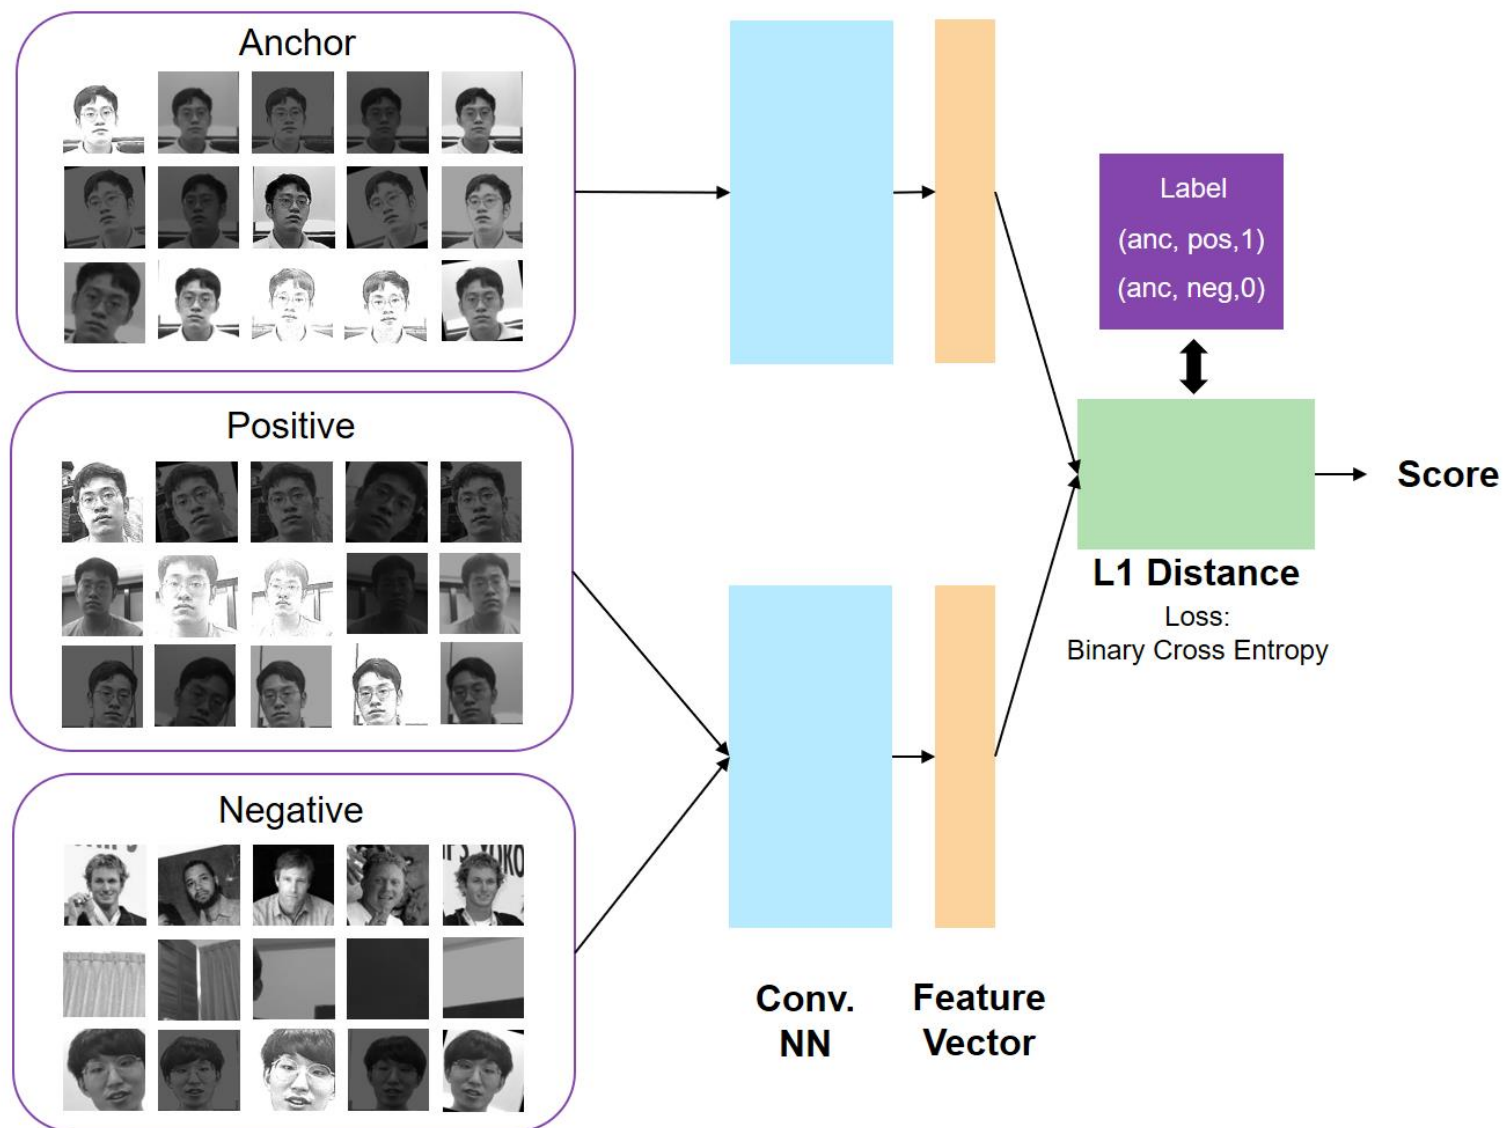
- Humans exhibit a strong ability to recognize new patterns
- A **few-shot learning** model



Design & Implementation

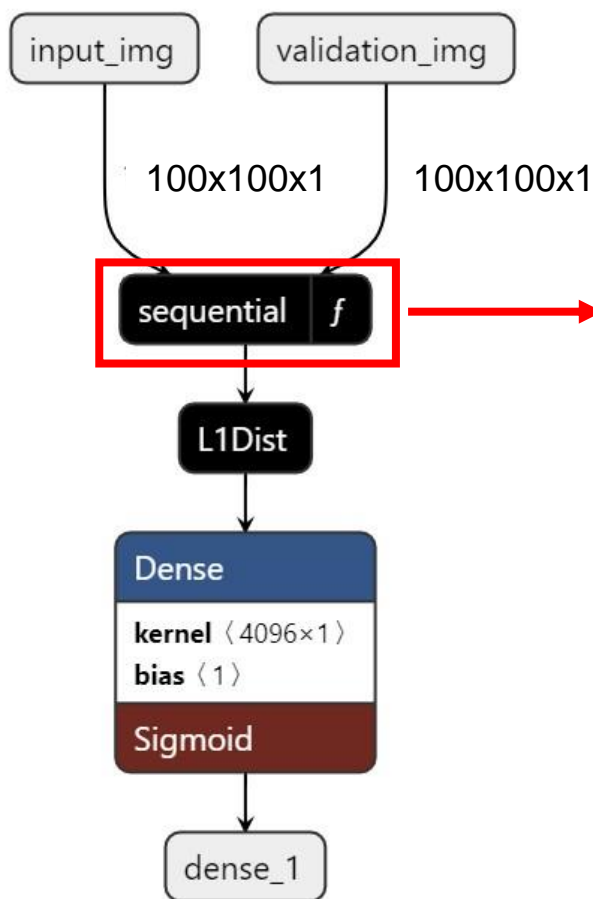
■ Data preprocessing

- **Resize: 100x100x1**
- **Data Augmentation**
- **Dataset**
 - Anchor: 600
 - Positive: 600
 - Negative: 600



Design & Implementation – Original Siamese Network

■ Inference on PC (Intel(R) Xeon(R) CPU E5-2620 v4 @ 2.10GHz)

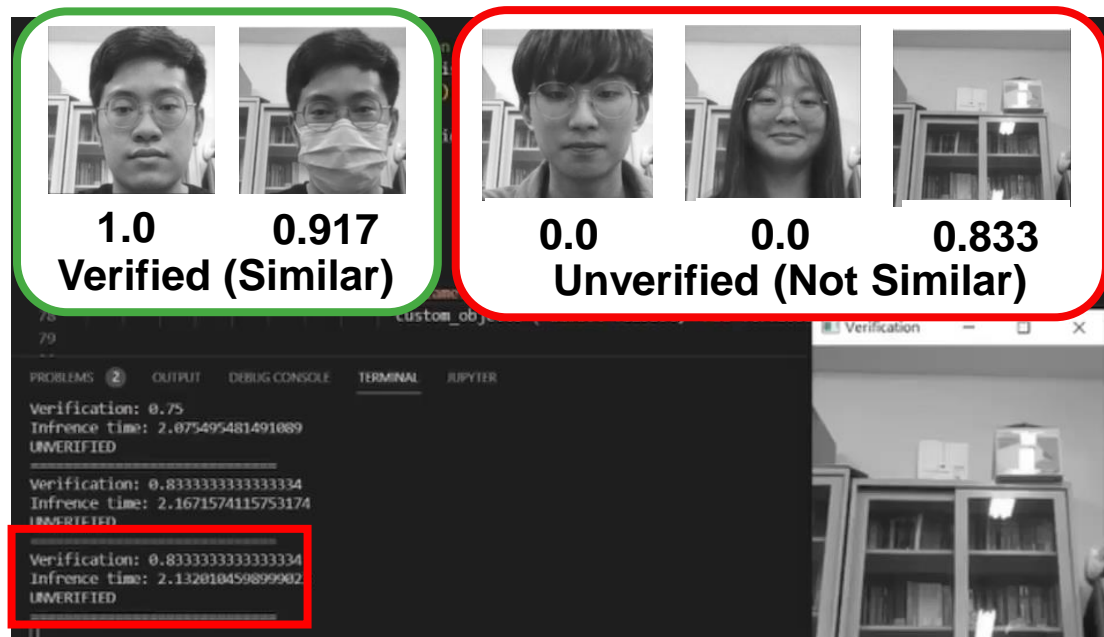



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 100, 100, 64)	6464
batch_normalization (Batch Normalization)	(None, 100, 100, 64)	256
activation (Activation)	(None, 100, 100, 64)	0
max_pooling2d (MaxPooling2D)	(None, 50, 50, 64)	0
conv2d_1 (Conv2D)	(None, 50, 50, 128)	401536
batch_normalization_1 (Batch Normalization)	(None, 50, 50, 128)	512
activation_1 (Activation)	(None, 50, 50, 128)	0
max_pooling2d_1 (MaxPooling2D)	(None, 25, 25, 128)	0
conv2d_2 (Conv2D)	(None, 25, 25, 128)	262272
batch_normalization_2 (Batch Normalization)	(None, 25, 25, 128)	512
activation_2 (Activation)	(None, 25, 25, 128)	0
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 128)	0
conv2d_3 (Conv2D)	(None, 12, 12, 256)	524544
batch_normalization_3 (Batch Normalization)	(None, 12, 12, 256)	1024
flatten (Flatten)	(None, 36864)	0
dense (Dense)	(None, 4096)	150999040
Total params: 152,196,160		
Trainable params: 152,195,008		
Non-trainable params: 1,152		

■ Result

- Float32
- Total parameters: **152.2 M**
- Accuracy: 0.967
- Avg. inference time: 2 s



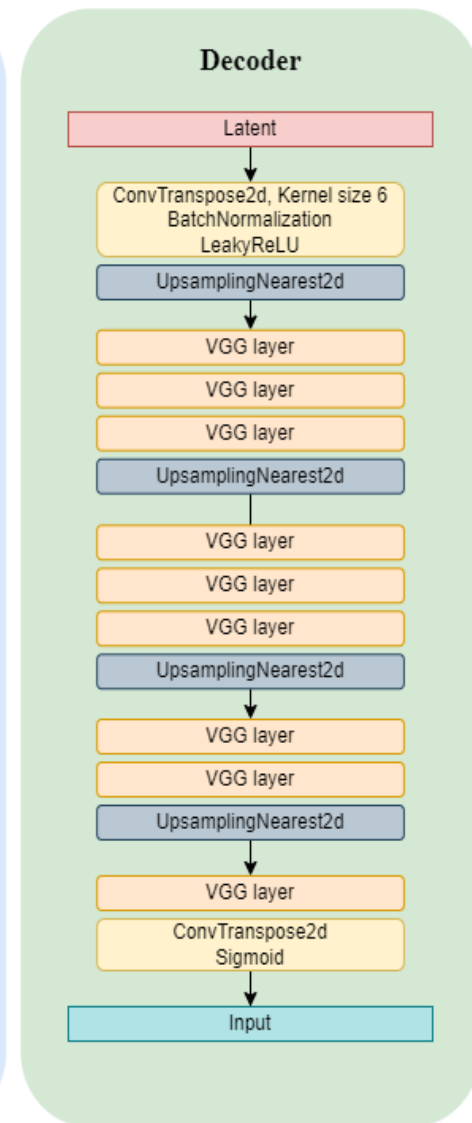
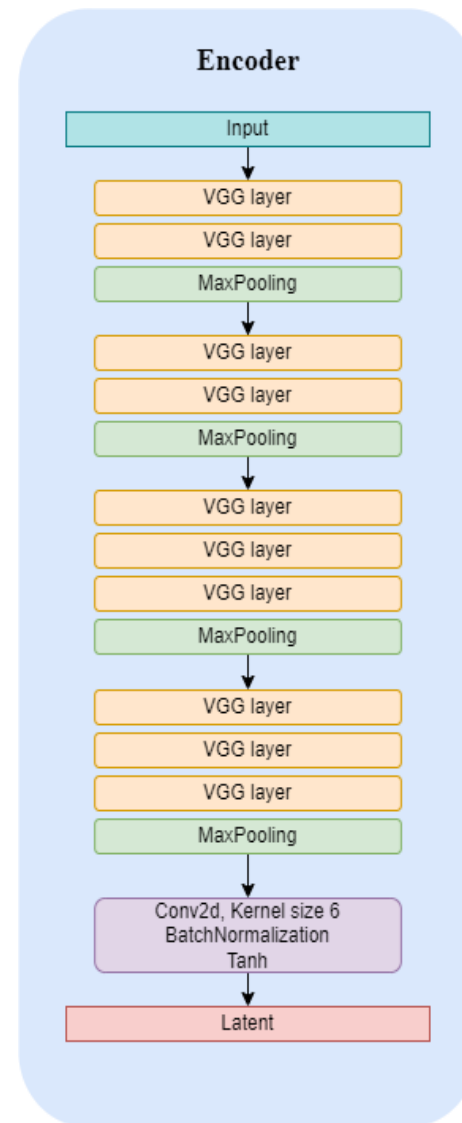

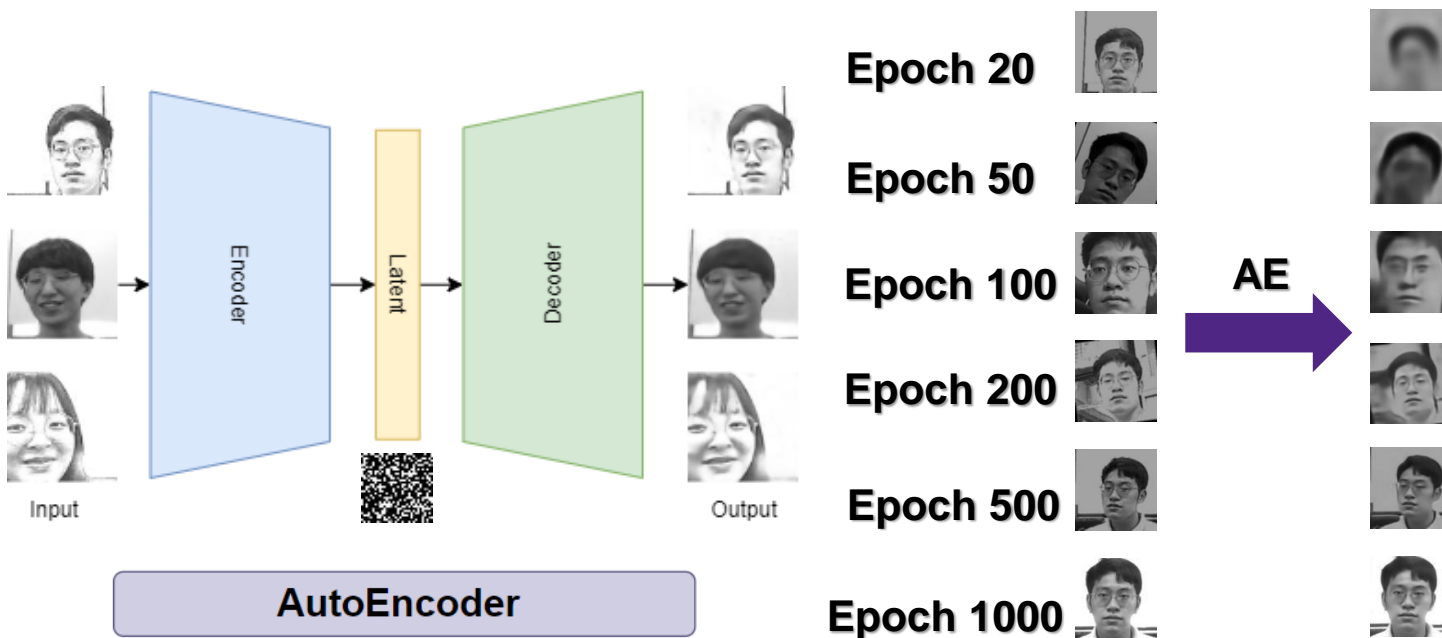
Design & Implementation – AutoEncoder

■ Dimension Reduction

- Compress input image into latent with Encoder
- Use latent to recover image with Decoder

■ Method

- Use VGG-16 as an architecture baseline

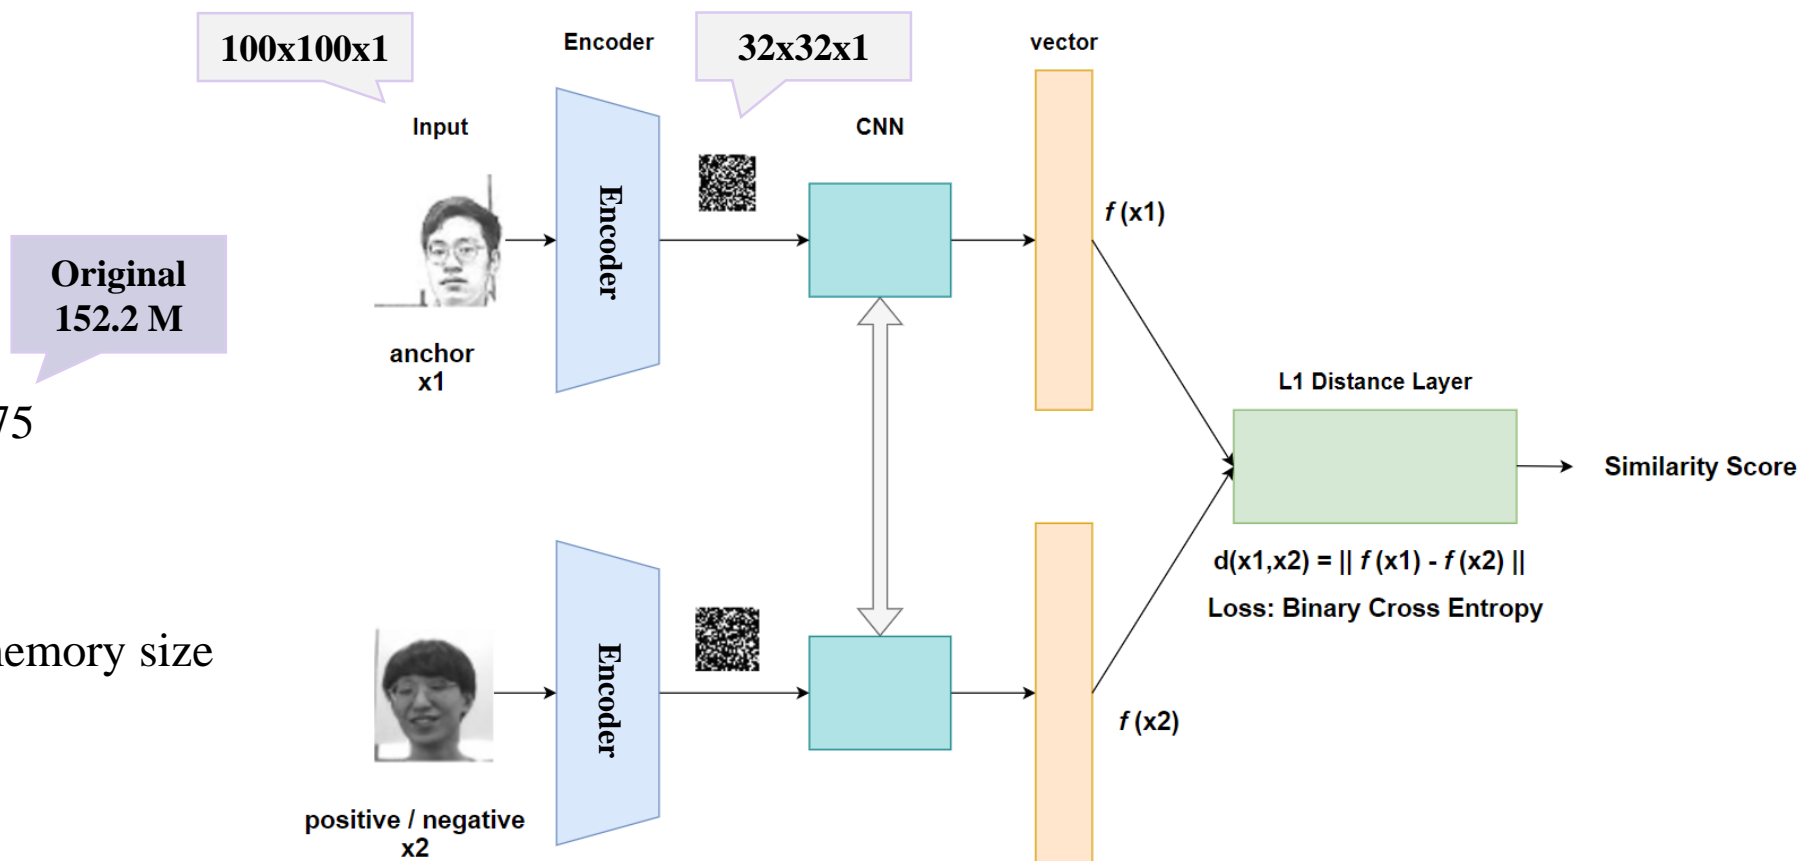


Design & Implementation – AutoEncoder + Siamese Network

■ Inference on PC (Intel(R) Xeon(R) CPU E5-2620 v4 @ 2.10GHz)

■ Result

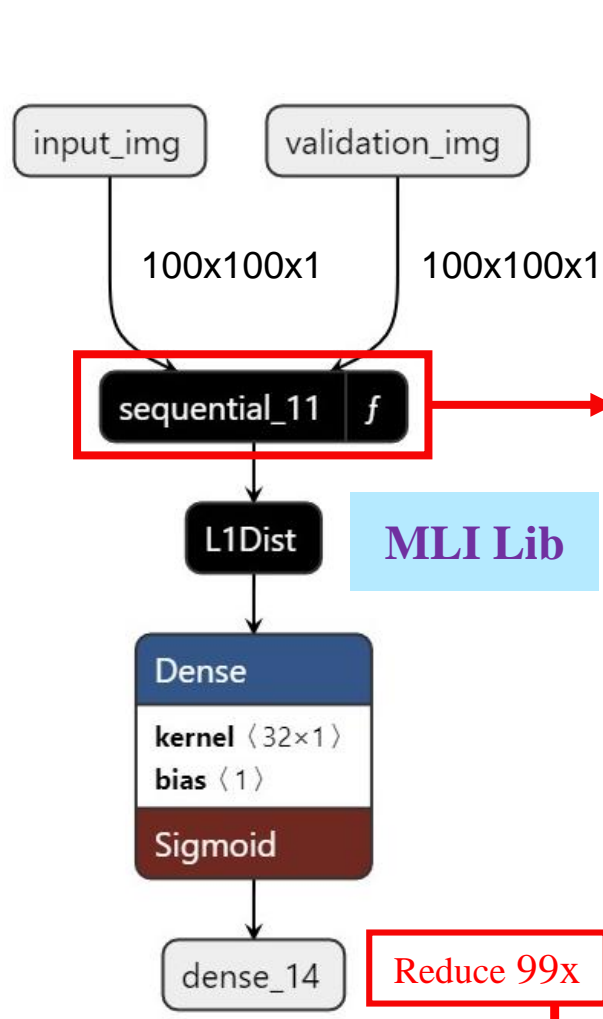
- Float32
- Parameters:
 - Encoder: 26,516,928
 - Siamese Network: 65,675
 - Total: **26.7 M**
- Memory Usage:
 - $\approx 500\times$ ARC EM9D memory size
- Accuracy: 0.98
- Inference time: 1.2 s



AutoEncoder + Siamese Network

Design & Implementation – Siamese Network (Fine tune)

■ Inference on ARC EM9D



MLI Lib

Reduce 99x

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 100, 100, 4)	40
depthwise_conv2d (DepthwiseC	(None, 100, 100, 4)	36
batch_normalization (BatchNo	(None, 100, 100, 4)	16
activation (Activation)	(None, 100, 100, 4)	0
max_pooling2d (MaxPooling2D)	(None, 50, 50, 4)	0
conv2d_1 (Conv2D)	(None, 50, 50, 8)	296
depthwise_conv2d_1 (Depthwis	(None, 50, 50, 8)	72
batch_normalization_1 (Batch	(None, 50, 50, 8)	32
activation_1 (Activation)	(None, 50, 50, 8)	0
max_pooling2d_1 (MaxPooling2	(None, 25, 25, 8)	0
conv2d_2 (Conv2D)	(None, 25, 25, 16)	1168
depthwise_conv2d_2 (Depthwis	(None, 25, 25, 16)	144
batch_normalization_2 (Batch	(None, 25, 25, 16)	64
activation_2 (Activation)	(None, 25, 25, 16)	0
max_pooling2d_2 (MaxPooling2	(None, 12, 12, 16)	0
depthwise_conv2d_3 (Depthwis	(None, 12, 12, 16)	144
batch_normalization_3 (Batch	(None, 12, 12, 16)	64
activation_3 (Activation)	(None, 12, 12, 16)	0
max_pooling2d_3 (MaxPooling2	(None, 6, 6, 16)	0
conv2d_3 (Conv2D)	(None, 3, 3, 32)	4640
depthwise_conv2d_4 (Depthwis	(None, 2, 2, 32)	288
batch_normalization_4 (Batch	(None, 2, 2, 32)	128
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 32)	4128
Total params: 11,260		
Trainable params: 11,108		
Non-trainable params: 152		

■ Description

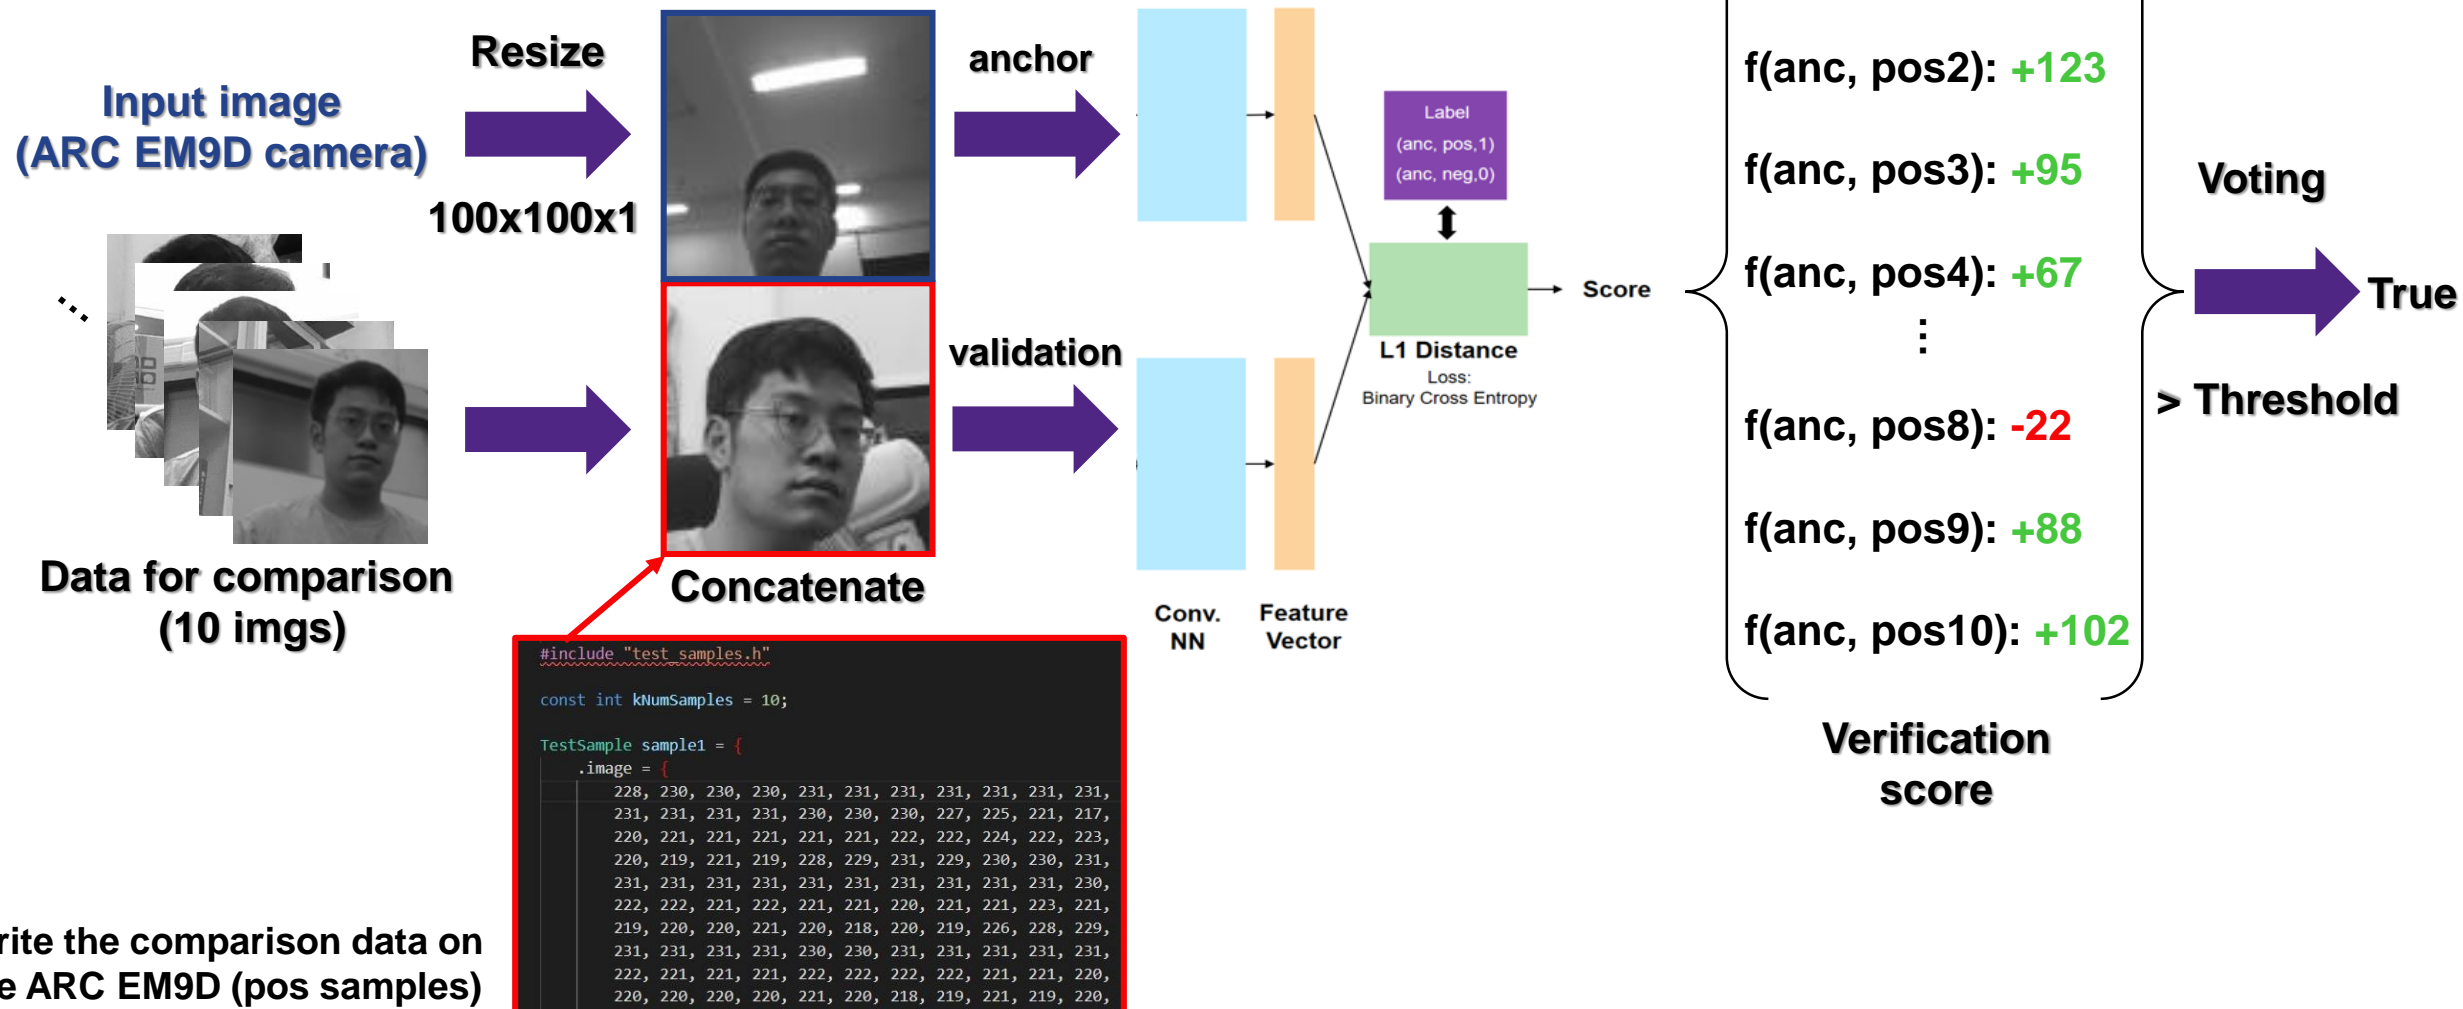
- Total parameters: **11,260**
- Int8 quantization
- Avg. inference time: 1 s

■ MLI Library (micro_op)

- AddConv2D()
- AddDepthwiseConv2D()
- AddMaxPooling2D()
- AddRelu()
- AddLogistic()
- AddFullyConnected()
- AddSub()
- AddMul()

Design & Implementation - Siamese Network (Fine tune)

■ Inference on ARC EM9D



Agenda

- Introduction
- Difficulties & Innovation
- Design & Implementation
- **Result & Demo**
- Q&A

Results & Demo – Siamese Network

	Original Siamese Network	AE + Siamese Network	Fine-tune (float32)	Fine-tune (int8)	Fine-tune (int8)
Processing Platform	*CPU	*CPU	*CPU	*CPU	ARC EM9D
Total Parameter	152.2 M	26.7 M	11,260	11,260	11,260
Can program on ARC EM9D ?	No	No	No	Yes (98%)	
Accuracy	0.967	0.98	0.933	0.8	N/A
Inference time	2.1 s	1.5 s	1.0 s	0.8 s	0.1 s

Agenda

- Introduction
- Difficulties & Innovation
- Design & Implementation
- Result & Demo
- **Q&A**

非接觸式控制面板 Contactless Control Panel

隊名：綠洲熊與他們的窩




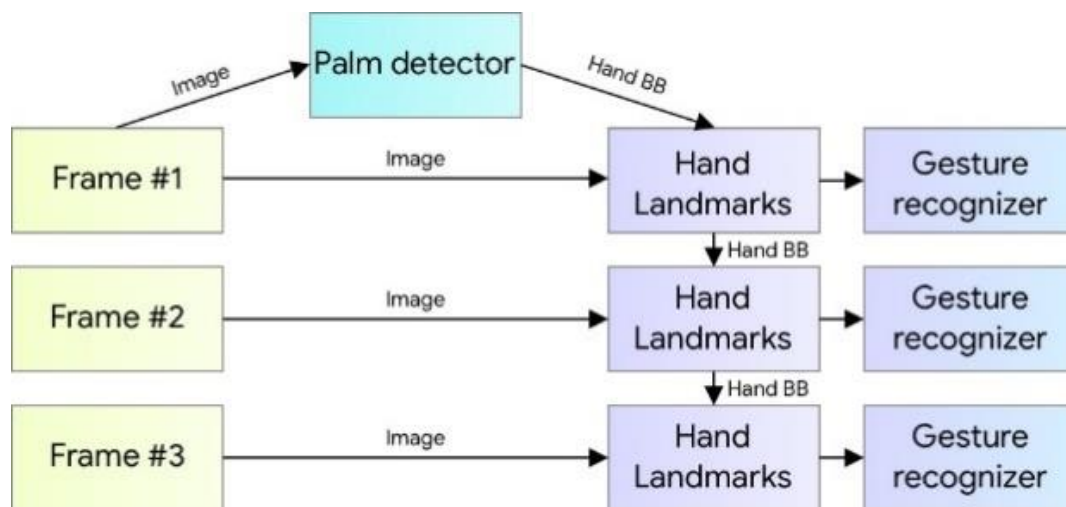
Appendix - Virtual Button Panel (NVIDIA Jetson Nano)

■ Hand Pose Estimation

- MediaPipe Hand

■ MediaPipe Hand Landmarks

- Palm detector model
- Hand landmarks model



- | | |
|-----------------------|-----------------------|
| 0. WRIST | 11. MIDDLE_FINGER_DIP |
| 1. THUMB_CMC | 12. MIDDLE_FINGER_TIP |
| 2. THUMB_MCP | 13. RING_FINGER_MCP |
| 3. THUMB_IP | 14. RING_FINGER_PIP |
| 4. THUMB_TIP | 15. RING_FINGER_DIP |
| 5. INDEX_FINGER_MCP | 16. RING_FINGER_TIP |
| 6. INDEX_FINGER_PIP | 17. PINKY_MCP |
| 7. INDEX_FINGER_DIP | 18. PINKY_PIP |
| 8. INDEX_FINGER_TIP | 19. PINKY_DIP |
| 9. MIDDLE_FINGER_MCP | 20. PINKY_TIP |
| 10. MIDDLE_FINGER_PIP | |

Appendix

■ Virtual Button Panel (NVIDIA Jetson Nano)



Select the virtual button




Click the virtual button




Implement on Jetson Nano

■ Simulation Animation (PC)




Move

Input Floor:
5
State:
MOVING
Going up




Open

Input Floor:
5
State:
OPENING
Going up




Close

Input Floor:
5
State:
CLOSING
Going up



Hold

Input Floor:
5
State:
HOLD
Going up



Multiple floors

Input Floor:
5 7 8 9
State:
MOVING
Going up

Reference

- [1] Koch, Gregory, Richard Zemel, and Ruslan Salakhutdinov. "Siamese neural networks for one-shot image recognition." ICML deep learning workshop. Vol. 2. 2015.
- [2] Sheng, T., Feng, C., Zhuo, S., Zhang, X., Shen, L., & Aleksic, M. (2018, March). A quantization-friendly separable convolution for mobilenets. In *2018 1st Workshop on Energy Efficient Machine Learning and Cognitive Computing for Embedded Applications (EMC2)* (pp. 14-18). IEEE.

Thank You

